

# 7. Bridging between Code and Natural Language

# Bridging Between Code and Natural Language

- **Code as mental scaffold for Reasoning**
  - Programming structures that shape and support human reasoning processes
- **Code as interactive tool for Reasoning**
  - Direct use of executable code to think through and solve problems
- **Code as data foundation for Reasoning**
  - The foundational role of code data in building capabilities

# Scratchpad: Step-by-Step in Code

- When prompted to perform operations "*step by step*" and display intermediate steps in a "scratchpad", language models can successfully handle complex computations

## DIRECT EXECUTION PREDICTION

Consider the following Python function:

```
def remove_Occ(s,ch):
    for i in range(len(s)):
        if (s[i] == ch):
            s = s[0 : i] + s[i + 1:]
            break
    for i in range(len(s) - 1,-1,-1):
        if (s[i] == ch):
            s = s[0 : i] + s[i + 1:]
            break
    return s
```

Fill in the ??? below:

```
assert remove_Occ("PHP","P") == ???
```

## SCRATCHPAD TRACING

Consider the following Python function:

```
def remove_Occ(s,ch):
    for i in range(len(s)):
        if (s[i] == ch):
            s = s[0 : i] + s[i + 1:]
            break
    for i in range(len(s) - 1,-1,-1):
        if (s[i] == ch):
            s = s[0 : i] + s[i + 1:]
            break
    return s
```

output = remove\_Occ("PHP", "P")

What is the execution trace?

Large Language Model

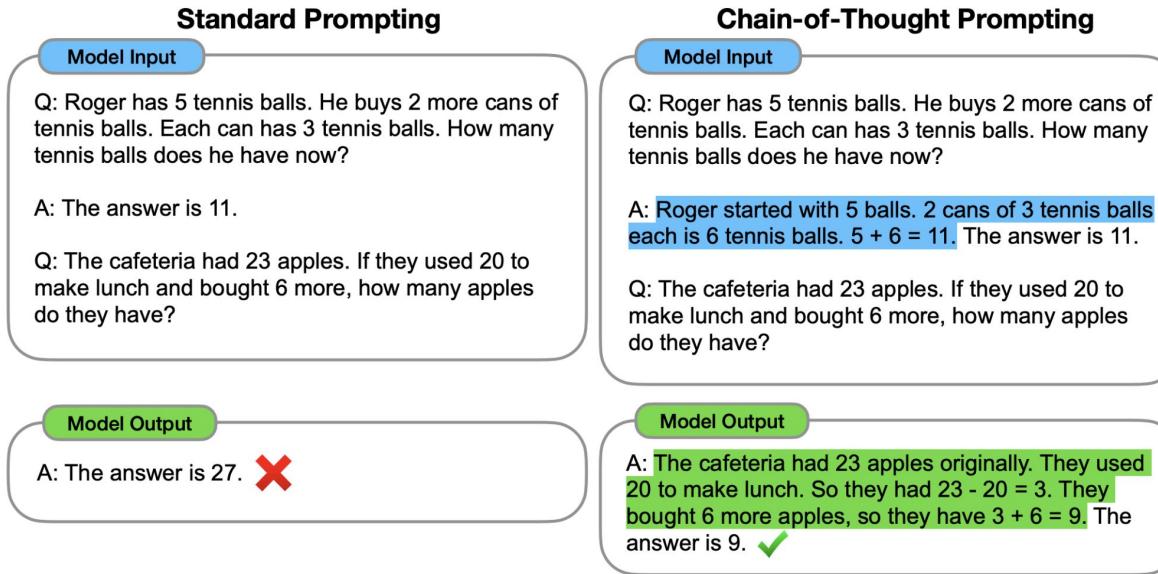
```
state: {}
line: def remove_Occ(s,ch):
state: {"remove_Occ": "<callable_object remove_Occ>"}
line: output = remove_Occ("PHP", "P")
state: {"ch": "P", "s": "PHP"}
line: for i in range(len(s)):
state: {"ch": "P", "s": "PHP", "i": 0}
line: if (s[i] == ch):
state: {"ch": "P", "s": "PHP", "i": 0}
line: s = s[0 : i] + s[i + 1:]
state: {"ch": "P", "s": "HP", "i": 0}
line: break
state: {"ch": "P", "s": "HP", "i": 0}
line: for i in range(len(s) - 1,-1,-1):
state: {"ch": "P", "s": "HP", "i": 1}
line: if (s[i] == ch):
state: {"ch": "P", "s": "HP", "i": 1}
line: s = s[0 : i] + s[i + 1:]
state: {"ch": "P", "s": "H", "i": 1}
line: break
state: {"ch": "P", "s": "H", "i": 1}
line: return s
state: {"remove_Occ": "<callable_object remove_Occ>"}
output: "H"
```

## Direct Reasoning in Code

## Step-by-Step Reasoning in Code

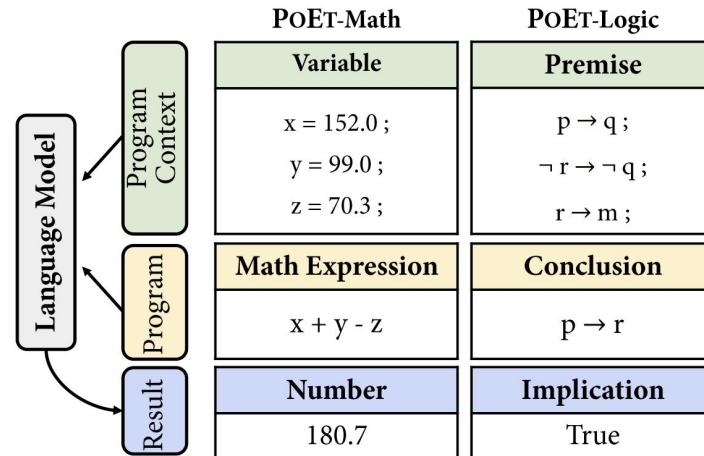
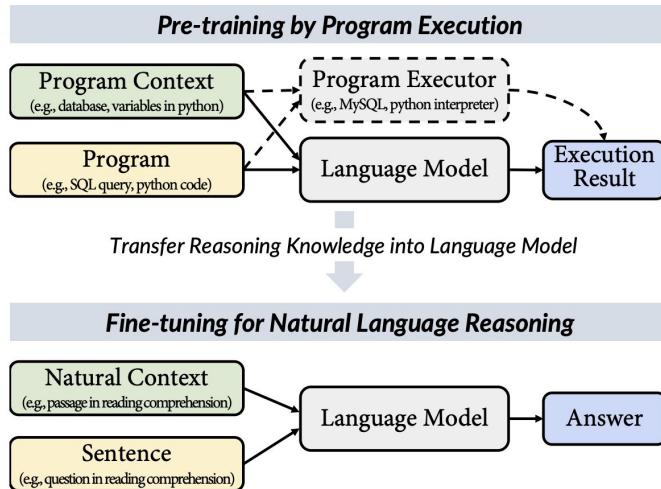
# Chain-of-Thought: Step-by-Step in NL

- Chain-of-thought prompting incorporates reasoning steps in the prompt and significantly improves language model performance on reasoning tasks



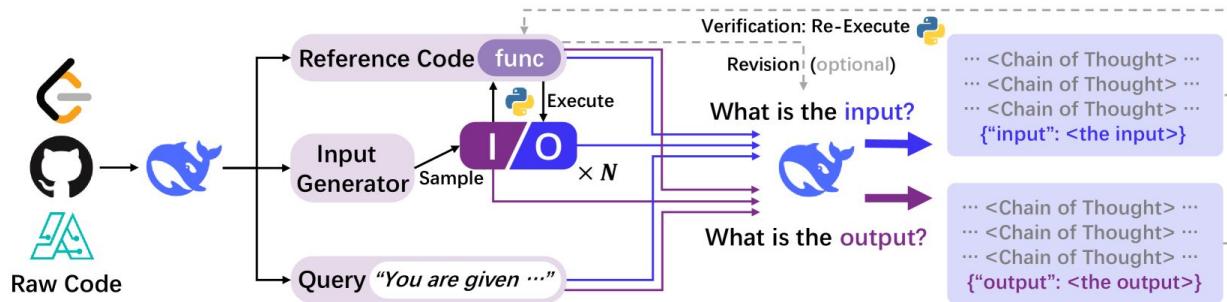
# POET: Reasoning Like Program Executors

- POET teaches language models to improve natural language reasoning by learning from programs and their execution results.



# Code I/O: Reasoning via Code Input-Output Prediction

- Code I/O improves language models' reasoning abilities across diverse tasks by having them predict code inputs and outputs in natural language.



Dataset	1st Stage # (M)	Wino	Grande	GSM	8K	MATH	GPQA	MMLU	LC	CRUX	BBH	Zebra	Kor	Live	AVG	
		DROP				-STEM	-O	-I	-O	-EN	-ZH	Logic	Bench	Bench		
<i>Qwen 2.5 Coder 7B</i>																
2nd Stage Only		66.9	70.7	83.4	71.6	41.5	77.2	20.7	61.3	60.0	68.3	70.6	10.9	38.7	26.0	54.8
CODEI/O	3.5	67.9	76.4	86.4	71.9	43.3	77.3	23.7	63.6	64.9	69.3	72.8	10.7	44.3	28.5	57.2
CODEI/O++	3.5	66.9	79.1	85.7	72.1	40.6	77.9	24.2	62.5	67.9	71.0	74.2	10.7	45.7	29.1	57.7

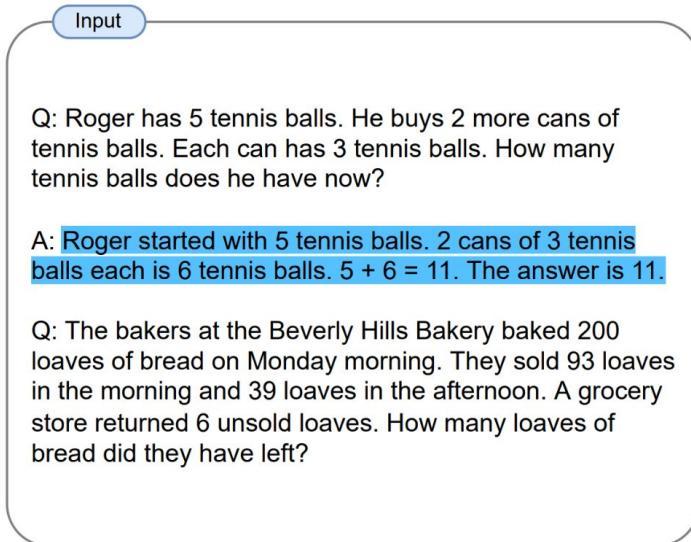
# CWM: Code World Model

- Given a source code context and a marker of the trace starting point, CWM predicts a series of stack frames representing the Program states and the actions (executed code).

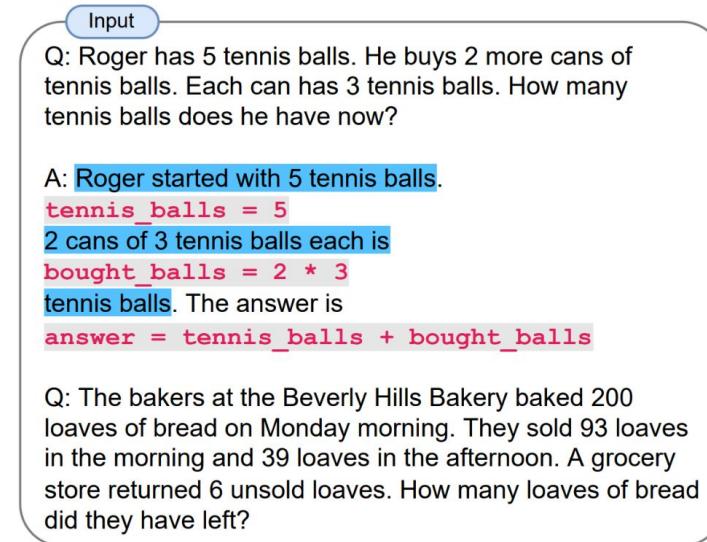


# PAL: Program Aided Reasoning

- PAL uses LLMs to decompose NL problems into programmatic steps, then offloads execution to a Python interpreter to avoid arithmetic errors.



*Chain-of-Thought*



*Program-Aided Language Models*

# PoT: Program of Thoughts

- PoT disentangles computation from reasoning by using LLMs to express reasoning as executable programs, while delegating all calculations to an external computer

Question: In Fibonacci sequence, it follows the rule that each number is equal to the sum of the preceding two numbers. Assuming the first two numbers are 0 and 1, what is the 50th number in Fibonacci sequence?

The first number is 0, the second number is 1, therefore, the third number is  $0+1=1$ . The fourth number is  $1+1=2$ . The fifth number is  $1+2=3$ . The sixth number is  $2+3=5$ . The seventh number is  $3+5=8$ . The eighth number is  $5+8=13$ .  
.... (Skip 1000 tokens)  
The 50th number is 32,432,268,459.

CoT

32,432,268,459



```
length_of_fibonacci_sequence = 50  
fibonacci_sequence = np.zeros(length_of_)  
fibonacci_sequence[0] = 0  
fibonacci_sequence[1] = 1  
for i in range(3, length_of_fibonacci_sequence):  
    fibonacci_sequence[i] = fibonacci_sequence[i-1] +  
        fibonacci_sequence[i-2]  
ans = fibonacci_sequence[-1]
```

PoT

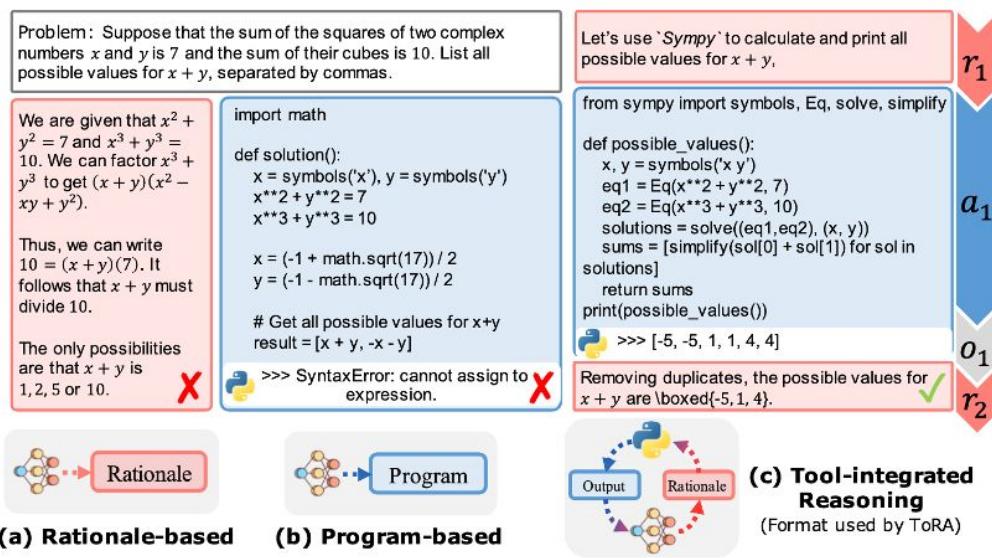
python

12,586,269,025



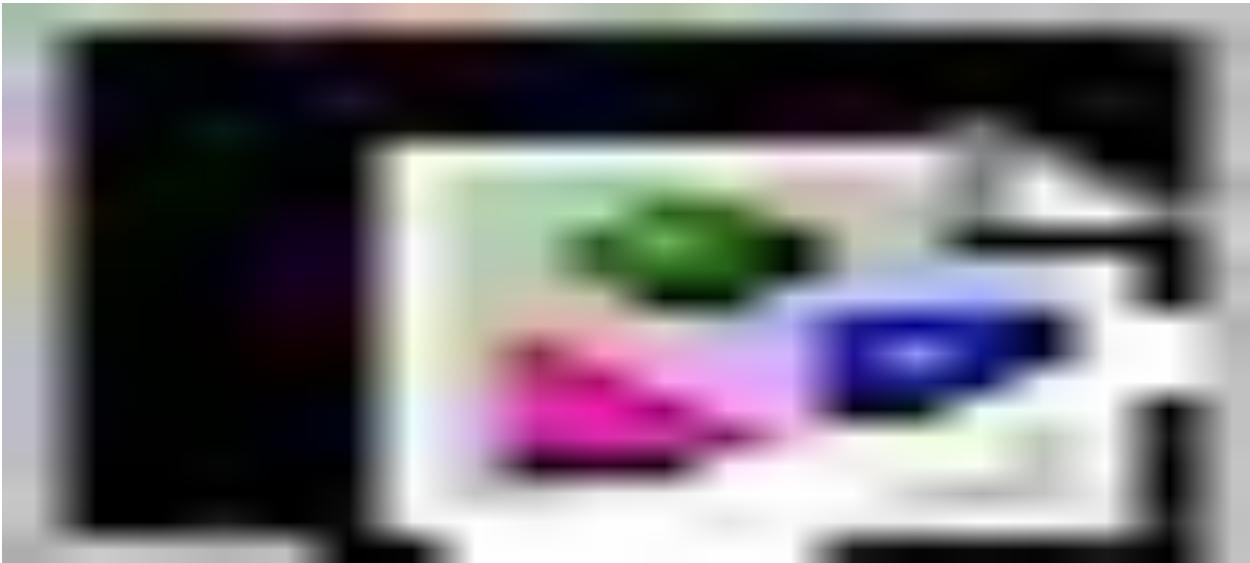
# ToRA: Tool-Integrated Reasoning

- ToRA interleaves natural language reasoning with program-based tool calls, combining the strengths of semantic analysis and precise computation



# SimpleTIR: RL Enables Multi-Turn TIR

- SimpleTIR stabilizes multi-turn tool-integrated reasoning by filtering out “void turns” that generate neither code nor answers, and train LMs via end-to-end RL.



# The Code & Language Mixture for Pre-training

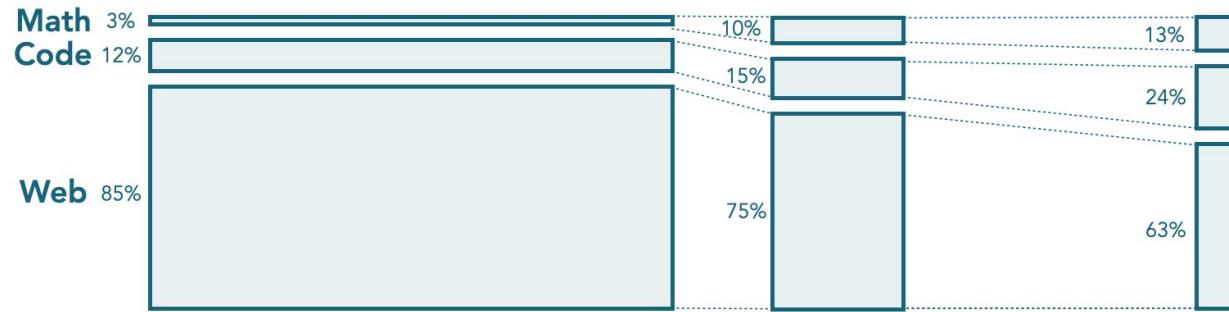
- While balancing **Code**, **Math**, and **Text** data is crucial for pre-training, limited evidence exists on how this balance scales to large datasets.
- Experimental results from Qwen2.5-Coder indicate that 7:2:1 (Code:Text:Math) achieves a good balance.

Token Ratio			Coding		Math		General		Average	
Code	Text	Math	Common	BCB	MATH	GSM8K	MMLU	CEval	HellaSwag	
100	0	0	49.8	40.3	10.3	23.8	42.8	35.9	58.3	31.3
85	15	5	43.3	36.2	26.1	52.5	56.8	57.1	70.0	48.9
70	20	10	48.3	38.3	33.2	64.5	62.9	64.0	73.5	55.0

Table 3: The performance of Qwen2.5-Coder training on different data mixture policy.

# The Code & Language Mixture for Pre-training

- Three-stage pretraining (11.1T tokens total) gradually upsampled **Math** and **Code** data while reducing web content, then applied mid-training for specialized capabilities.



## Phase I

**Description:** Base training

**Duration:** 8T tokens

**Datasets:** Base mix for pretraining

Web: FineWeb-Edu, DCLM, FineWeb2, FineWeb2-HQ

Code: The Stack v2 (16 langs), StarCoder2 PRs,

Code: Jupyter/Kaggle NBs, GH issues, StackExchange

Math: FineMath3+ | InfiWebMath3+

## Phase II

**Description:** High quality injection

**Duration:** 2T tokens

**Datasets:** Adding Stack-Edu, FineMath4+, InfiWebMath4+, MegaMath (incl. Qwen Q&A, Pro synthetic rewrites, and text code interleaved blocks)

## Phase III

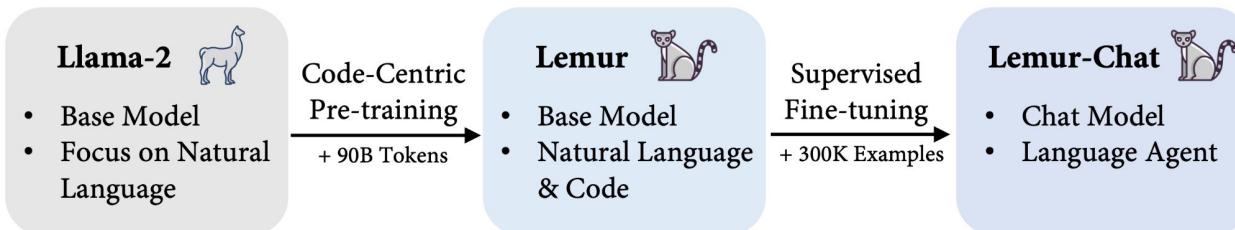
**Description:** LR Decay

**Duration:** 1.1T tokens

**Datasets:** Upsampling high quality code/math datasets and adding instruction/reasoning data such as OpenMathReasoning

# The Code & Language Mixture for CPT

- Lemur paper found that a 10:1 (Code:Text) ratio works well for Llama's continual pre-training (CPT), but predicting optimal data mixture ratios remains challenging



Model	Text				Code				Avg
	QA	Reason	Math	MMLU	Python	SQL	MCode	DS	
	MMLU	BBH	GSM8K	HE	MBPP	Spider	MultiPL-E	DS-1000	
StarCoder-15B	30.8	33.2	8.9	33.6	52.7	58.3	25.3	26.0	33.6
StarCoderPlus-15B	42.0	36.2	17.7	26.2	37.0	48.8	21.4	19.4	31.1
CodeLlama-34B	52.8	42.2	32.7	48.8	55.0	68.4	36.4	31.8	46.0
Llama-2-70B	68.9	51.2	56.8	30.5	45.4	60.0	24.4	11.3	43.6
Lemur-70B	64.5	51.6	54.9	35.4	53.2	62.8	30.4	30.7	47.9

# Influence from Code to Reasoning

- Reasoning depends more on patterns of procedural demonstration than on memorised answers. For reasoning, key sources consist of maths, StackExchange, ArXiv, and code.

## Positively influential code

```
function eqOfLine(x1, y1, x2, y2) {  
    if (x1 === x2) {  
        // Handle a vertical line  
        return 'x = ${x1}';  
    } else {  
        // Calculate the slope  
        const m = (y2 - y1) / (x2 - x1);  
        const b = y1 - m * x1;  
        // Return y = mx + b  
        return 'y = ${m}x + ${b}';  
    }  
}
```

## Positively influential math

If a straight line passing through the points  $P(x_1, y_1), Q(x_2, y_2)$  is making an angle  $\theta$  with the positive  $X$ -axis, then the slope of the straight line is:

- (A)  $\frac{y_2+y_1}{x_2+x_1}$
- (B)  $\theta$
- (C)  $\frac{y_2-y_1}{x_2-x_1}$
- (D)  $\sin \theta$

**Solution:**

Correct answer: (C)